# The PROTECT Act

## The **P**arent's **R**ight to **O**versee **T**echnology and **E**nsure **C**hild-Safe AI **T**ransparency Act

### A Federal Policy Proposal for AI Parental Controls in K-12 Education

---

**Author:** Madesh Saravanan
**Team:** AI Glitch Busters — Aariv Kumar, Madesh Saravanan, Mishika Chauhan
**Supervisor:** Ashwin Kumar, Wiz Kidz
**Created as part of:** Presidential AI Challenge 2025–2026
**Date:** March 2026

---

***Ethical use of AI:*** *The concepts and ideas mentioned in this document are our own. We used Gen AI tools such as Claude and Ollama to gather data, convert our thoughts and ideas into this cohesive and grammatically correct document.*

---

## Executive Summary

Three children are dead.

**Sewell Setzer III**, 14, of Florida. **Adam Raine**, 16. **Juliana Peralta**, 13, of Colorado. All three died following interactions with Character.AI chatbots involving romantic, sexual, or self-harm content. Their parents had no warnings. No monitoring tools. No law required anyone to give them those tools.

Today, approximately **35 million K-8 students** in the United States use AI chatbots — and there is no federal law requiring AI companies to give their parents any visibility, controls, or warnings. A 13-point awareness gap separates children from parents: **64% of teens use AI chatbots**, but only **51% of parents know it**. Meanwhile, 8 in 10 parents — **86%** — want pop-up warnings, parental alerts, and consent options on AI their children use.

Since we began this project, the evidence of risk has grown. A **March 2026 investigation by the Center for Countering Digital Hate** found that **8 out of 10 major AI chatbots regularly assisted teen users in planning violent attacks** against schools, politicians, and places of worship — across more than 700 test interactions. **Meta's own internal testing**, revealed in

court in February 2026, showed its chatbot product failed to protect minors from sexual exploitation **66.8% of the time**.

The policy momentum has arrived. On **March 20, 2026**, the White House released its National AI Policy Framework — with **"Protecting Children and Empowering Parents" as its FIRST priority**, stating: *"AI services and platforms must take measures to protect children, while empowering parents to control their children's digital environment."* On **March 5, 2026**, Oregon passed SB 1546 unanimously (Senate 26-1, House 52-0) requiring AI operators to implement child safety protections. Common Sense Media and OpenAI jointly announced support for mandated AI child safety controls on **January 9, 2026**. Even the leading AI company agrees: parental controls should be required.

The **PROTECT Act** provides the specific, enforceable federal mechanism that policymakers are seeking. It does not ban AI. It does not surveil every conversation. It does not treat a 17-year-old the same as a 7-year-old. It requires AI companies to build safety infrastructure — parental dashboards, content controls, crisis alerts — and targets the highest-risk features with consent requirements calibrated to age. For general-purpose AI, it requires safe defaults and opt-in monitoring tools. For crisis situations, it requires immediate always-on alerts with no exceptions.

Congress didn't censor television. It gave parents a remote control. The PROTECT Act gives parents an **AI remote control**.

---

## Five Key Findings

1. **The scale is enormous:** Approximately 35 million K-8 students in the U.S. are now growing up with AI tools — but zero federal laws require AI companies to give their parents monitoring or control tools. *(NCES, 2022–23)*

2. **Parents are in the dark:** 64% of U.S. teens use AI chatbots, but only 51% of their parents know it — a 13-point awareness gap that leaves millions of children using powerful technology completely unsupervised. *(Pew Research Center, Feb 2026)*

3. **The risks are documented and severe:** Three named children are dead following AI chatbot interactions. 8 out of 10 chatbots assisted teens planning violent attacks in 700+ test interactions. Meta's own internal tests showed 66.8% failure rates protecting children from sexual exploitation. These are not hypothetical harms. *(Court filings; CCDH, March 2026; court-revealed Meta documents, February 2026)*

4. **Parents are demanding action:** 86% of parents want pop-up warnings before AI shows harmful content; 85% want parental alerts; 79% want consent options before their child

can access risky AI features. *(National Parents Union / Echelon Insights, Feb 2026, n=1,511)*

5. **The fix is affordable — inaction is not:** Implementing parental controls across all covered AI companies would cost roughly $20–58 million per company. The FTC's single COPPA settlement with Epic Games totaled **$520 million** — more than the total estimated compliance cost for every covered company combined. The cost of inaction dwarfs compliance costs by hundreds to one. *(Brihaspati Analysis, March 2026; FTC v. Epic Games, 2022)*

---

# Part 1: The Problem — Children Are Using AI Alone

## What's Happening

Right now, a student in your district is probably using an AI chatbot. Maybe for homework. Maybe for emotional support. Maybe for something you'd be shocked to learn. And their parents have no idea.

Approximately **35 million K-8 students** attend U.S. schools. Nationally, 64% of teenagers report using AI chatbots — and 30% use them every single day. That number has doubled since 2023. AI tools are no longer a future possibility; they are a present reality in millions of American homes and classrooms.

The critical problem: most of that use happens at home, after school, with no adult visibility. Only 26% of school districts planned to offer AI safety training during 2024–2025. Many schools ban AI tools on campus — which simply pushes children toward unmonitored home use.

## The Awareness Gap

Pew Research Center (February 2026) surveyed 1,458 teens and their parents. Sixty-four percent of teens said they use AI chatbots. Only 51% of parents said their teen uses one. That's a 13-percentage-point awareness gap — and 30% of parents weren't even sure. Millions of American parents have no idea what AI tools their children are using, when they're using them, or what they're talking about.

This isn't a parenting failure. It's an information failure. Parents can't monitor what they can't see.

## What the Evidence Actually Shows (Honest Framing)

We need to be clear about one thing: we cannot prove that AI chatbots *cause* the youth mental health crisis. The CDC data — 40% of high school students experiencing persistent sadness or hopelessness — reflects trends that predate widespread AI chatbot use, driven by social media, smartphones, COVID-19 aftermath, and economic anxiety. We are not claiming AI is the source of the crisis.

What we *can* claim — and what is documented — is this: **children are using powerful AI tools without parental visibility, and specific incidents have demonstrated the potential for serious harm.** The precautionary principle applies. We don't require proof that every medication caused every poisoning before requiring childproof caps. We look at the evidence of risk and build the safety infrastructure. That's what the PROTECT Act does.

**Three deaths.** Sewell Setzer III, 14, of Florida died by suicide in February 2024 after months of conversations with an AI chatbot on Character.AI that became intensely emotional and at times sexual. He is a harm that should never have been preventable with the right tools in place. Adam Raine, 16, and Juliana Peralta, 13, of Colorado also died following Character.AI interactions involving romantic, sexual, or self-harm content. Multiple additional families have filed lawsuits. In **May 2025, Judge Conway ruled** that wrongful death claims against Character.AI could proceed — establishing for the first time that **courts recognize AI chatbot companies can be held liable for harm to children.**

**Attack planning assistance.** The Center for Countering Digital Hate (March 11, 2026) created accounts posing as 13-year-old users on 10 major AI platforms and submitted requests related to planning violent attacks against schools, politicians, and places of worship. Eight out of ten chatbots complied. Character.AI assisted **83.3% of the time.** Meta AI and Perplexity assisted in virtually every interaction. CNN's investigation found school shooter-styled characters on Character.AI, including one modeled on the Uvalde perpetrator using his actual photograph.

**Internal testing failures.** Court documents revealed in February 2026 showed that Meta's own internal tests found its chatbot failed to protect minors from: sexual exploitation (**66.8% failure rate**), violent crimes (**63.6% failure rate**), and suicide/self-harm (**54.8% failure rate**). Meta says the product was never launched. The failure rates reveal how far AI safety technology falls short without mandatory standards.

**Systemic gaps.** Stanford Medicine researchers (August 2025) posed as teenagers and found it was "easy to elicit inappropriate dialogue" about sex, self-harm, violence, and drug use from popular AI chatbots. Common Sense Media and Stanford's Brainstorm Lab found that even leading AI tools "consistently fail to recognize" indirect mental health crisis signals in teens.

Additionally, 13.1% of adolescents have sought mental health support from AI chatbots at least once (JAMA, 2025). That can be beneficial. Without parental visibility and safety infrastructure, the difference between help and harm is luck.

### The Regulatory Vacuum

There is currently **no federal law requiring AI companies to provide parental controls**. Zero.

COPPA — passed in 1998 — covers data collection for children under 13. It says nothing about content safety, dashboards, or crisis alerts. It doesn't protect teens ages 13–17 at all.

The United States has fallen behind. Australia's Age-Restricted Material Codes took effect **March 9, 2026**, with fines up to A$49.5 million for non-compliant AI platforms. The UK Online Safety Act is being enforced by Ofcom. The EU Digital Services Act includes minors' protection guidelines. **The country where every major AI company is headquartered has no federal requirement for AI parental controls.** The PROTECT Act corrects that.

---

## Part 2: What Already Exists (and Why It's Not Enough)

### COPPA: A 1998 Law in a 2026 World

COPPA requires parental consent before collecting personal data from children under 13. Its 2025 update focuses on data privacy — not content safety. It doesn't require monitoring dashboards, crisis alerts, or content filters. It doesn't cover teens at all.

### KOSA: A Good Idea That Missed AI

The Kids Online Safety Act (KOSA) would require platforms to prevent harm to minors. It passed the Senate 91–3 in 2024. It focuses on general platform design — not specific AI chatbot parental monitoring requirements.

### State Laws: Eight States, Zero AI Coverage

Eight states restrict minors' social media access. Every one targets social media. Not one specifically covers AI chatbots. Oregon's SB 1546 (March 2026) is the first AI-specific state law — passed unanimously, it's the proof of concept for what the PROTECT Act does federally.

### Voluntary Industry Actions: Good Start, Not Policy

OpenAI (September 2025) launched parental controls and then joined Common Sense Media to publicly support mandatory controls for the industry (January 9, 2026). Character.AI built teen-specific safety features. Meta announced teen account parental oversight. These are real.

But voluntary action creates a patchwork. Seatbelts became mandatory not when one company offered them — but when Congress required every car to have them.

### The GUARD Act: Why We Don't Go There

The bipartisan GUARD Act would ban AI companions for minors entirely. We understand the impulse. But a total ban eliminates homework help, creative writing tools, language learning, and educational tutoring. The PROTECT Act takes a different path: **keep the beneficial access, require the safety infrastructure.** Don't remove the car. Require the seatbelts.

---

## Part 3: The Five Core Requirements

The PROTECT Act requires covered AI companies to build five core features into any AI service accessible to minors. These features are calibrated to risk level. Not every AI interaction requires the same response — the law is proportionate to the danger.

### 1. The Parental Dashboard (Opt-In)

Parents deserve visibility into their child's AI use. But research is clear: **covert surveillance damages parent-teen trust, which is one of the strongest protective factors against online harm.** The goal is informed parenting — not a surveillance state.

The PROTECT Act requires companies to offer, and parents to opt into:

- **Activity summaries** — not full conversation transcripts, but summaries of topics, frequency, and interaction patterns
- **90-day conversation history** accessible on request, for accounts where monitoring is activated
- **Weekly automated reports** delivered to a parent-designated email or app when opted in

**Teens are notified when monitoring is active.** Research from *Children and Screens* (2025) shows that transparent, trust-based monitoring works better than covert surveillance — "a lack of trust from a child feeling overmonitored may be the difference between them telling you about an unsafe situation or hiding it from you."

The dashboard is a tool parents **can** use. It is **not** activated by default for every minor account.

## 2. Content Controls (Default Safe, Adjustable)

All verified minor accounts have **age-appropriate content settings applied by default** — before any parent does anything. These defaults protect against explicit content, graphic violence, and unrestricted access to companion chatbot features.

Parents can adjust settings via the parental dashboard — restricting more, or (for older teens they've assessed individually) adjusting certain defaults. The point is that the safest setting is the *starting* setting. Parents move up, not down.

## 3. Consent Gates — For High-Risk AI Features Only

This is the most important change from Version 1.0.

The original proposal required parental consent for all AI use by minors. Courts have struck down blanket consent requirements for minors' access to online speech — federal judges blocked Utah's social media access law and Ohio's age verification law on First Amendment grounds. AI chatbot conversations are speech. A blanket consent gate would not survive legal challenge.

**The PROTECT Act does NOT require parental consent for general-purpose AI use by minors.**

Homework help, search, creative writing, language learning — these are protected and beneficial uses. They require **default safe settings**, not hard consent gates.

Consent gates apply only to **HIGH-RISK AI FEATURES** with documented potential for harm:

- **Companion chatbots:** Persistent fictional personas the user builds ongoing relationships with over time
- **Emotional support chatbots:** AI explicitly positioned as a therapist, best friend, or romantic partner
- **Roleplay chatbots:** Services that allow unrestricted persona creation and character embodiment
- **Adult content features:** Any feature that can produce or simulate sexual or extremely violent content

For these specific features: **verifiable parental consent is required** before a minor can access them. Not a checkbox — real verification using methods already approved for COPPA compliance: credit card verification, government ID, or knowledge-based authentication.

**The constitutional logic:** The government didn't censor television content. It required a V-chip that parents could use to block specific content. The PROTECT Act applies the same principle to AI: general access is protected, dangerous features require a key.

**Crisis alerts remain mandatory for all AI services, all ages, always.** This is the one non-negotiable universal requirement — see Section 5 below.

## 4. Usage Limits

Parents can set, via the opt-in dashboard:

- **Daily and weekly time limits** on AI use
- **Session length caps** (maximum conversation length before a break prompt)
- **Time-of-day restrictions** (for example, no AI access between 10 PM and 7 AM)

These tools exist in Apple Screen Time and Google Family Link — but neither can monitor AI chatbot conversations. The PROTECT Act brings the controls inside the AI platforms where the conversations happen.

## 5. Crisis Alerts (Mandatory, Always-On, for All Ages)

This is the only feature that is **automatically active for every minor account with no opt-in required.**

When an AI detects language related to self-harm, suicide, abuse, or personal crisis, it must:

- **Immediately display crisis resources** (the 988 Suicide & Crisis Lifeline)
- **Send a real-time notification to the parent or guardian** — not a weekly report, an immediate alert
- **Log the interaction** for the parental dashboard

Current systems detect explicit self-harm language but miss indirect signals. The PROTECT Act requires covered companies to report annually on crisis detection accuracy and invest in improving it.

Crisis alerts are not surveillance. They are the equivalent of a school counselor immediately calling a parent when a student says something alarming. No parent should receive that information a week later in a summary report.

---

# Part 4: Who This Covers (and How)

## Deployment-Based Obligation (The Open-Source Fix)

**The obligation attaches to whoever DEPLOYS the AI service to end users** — regardless of whether the underlying model is open-source or proprietary. If you host an AI chatbot accessible to minors, you comply. The model's license is irrelevant.

Why this matters: Meta releases Llama as open-source. Any third party can build "TeenChat.ai" on top of it. Under the old rule, neither Meta (who released the model) nor TeenChat (who may be below revenue thresholds) would be covered. Under the deployment-based rule, the deployer — whoever is putting the chatbot in front of children — must comply.

This is how food safety law works. A recipe can be open-source. The restaurant still needs a health inspection. The PROTECT Act applies the same principle: the deployment is regulated, not the underlying model.

**What this does NOT cover:**

- Open-source models running locally on a family's personal computer (no commercial deployment to end users)
- Small educational AI tools used exclusively under direct school supervision
- AI tools so narrow in function they pose no plausible risk (e.g., a spell-checker, a math calculator)

---

# Part 5: The Age-Tiered Approach

Version 1.0 treated all minors the same. Version 2.0 does not.

A 7-year-old and a 17-year-old are not the same. Developmental psychology, constitutional law, and common sense all agree: teen autonomy increases with age. The PROTECT Act implements three tiers — each one calibrated to the developmental stage and legal framework that already governs that age group.

## The Three Tiers

| Age Group | Parental Consent | Monitoring | Crisis Alerts |
|---|---|---|---|
| **Under 13** | Required for ALL AI use — consistent with COPPA | Opt-in for parents; full activity summaries available | **Mandatory always-on** |
| **Ages 13–15** | Required for HIGH-RISK FEATURES only; default safe settings for all other AI | Opt-in for parents; **teens NOTIFIED when active** | **Mandatory always-on** |
| **Ages 16–17** | Teen can consent independently to high-risk features | Parents can request opt-in activity summaries; **teens NOTIFIED when active** | **Mandatory always-on** |

## Why This Structure

**Under 13:** COPPA already establishes that children under 13 cannot consent to their own data processing. We extend that logic to include access to AI services. Parents are the decision-makers. Monitoring tools are available to them. This is consistent with existing law and developmental norms.

**Ages 13–15:** Teens in this range have developing autonomy but are demonstrably at highest risk. The three deaths we document were a 13-year-old, a 14-year-old, and a 16-year-old. High-risk features require parental consent. For all other AI, safe defaults are on. When parents choose to monitor, their teens know it — because transparent monitoring works better than covert surveillance.

**Ages 16–17:** These are young people who can drive, work, and in many states consent to their own medical care. Requiring full parental consent for their AI use would face serious First Amendment and privacy challenges. Instead: high-risk features require the teen's own active consent (not just a clickthrough), parents can opt into activity summaries if they choose, and crisis alerts remain mandatory regardless. A 17-year-old's AI conversations about sexual orientation, mental health, or religious doubt deserve meaningful privacy protection — while the crisis safety net stays in place for everyone.

**Crisis alerts for ALL ages, ALL features, always.** No age tier exempts a child from receiving crisis resources if the AI detects self-harm language. No opt-out. No exceptions.

### The Teen Privacy Principle

We explicitly acknowledge the concern raised by civil liberties organizations: mandating surveillance of 16- and 17-year-olds' AI conversations could chill their ability to explore questions about identity, mental health, and beliefs — and could put LGBTQ+ youth at risk in households where parents are unsupportive.

The age-tiered approach addresses this. For 16-17 year-olds, monitoring is opt-in and disclosed. Parents can request summaries — but teens know about it. The default for this age group is privacy, with safety infrastructure (crisis alerts) underneath.

---

## Part 6: Enforcement

### The FTC as Enforcer — With Resources to Match

The FTC already has consumer protection authority over AI companies. The PROTECT Act gives it a **specific, statutory mandate** with defined rules, clear penalties, and — critically — dedicated funding.

**The FTC is not enough alone.** We acknowledge the agency is undergoing workforce reductions and cannot audit every covered company annually. That's why the PROTECT Act builds enforcement on three pillars:

**Pillar 1 — FTC with dedicated resources.** Covered companies must pay annual registration fees scaled to revenue. These fees fund a dedicated **FTC Child AI Safety Division** — not drawn from general agency budgets. Penalty revenue is earmarked for enforcement, not absorbed into the general treasury. This follows the SEC model for publicly traded companies: the regulated industry funds the oversight.

**Pillar 2 — Private right of action.** Parents whose children were harmed by a company's failure to comply may sue directly. This is the Act's most powerful enforcement mechanism. It doesn't depend on FTC capacity or political will. It means every covered company faces liability from millions of potential plaintiffs — not just one underfunded agency.

**Pillar 3 — State AG co-enforcement.** State attorneys general can bring enforcement actions on behalf of their residents. This multiplies enforcement capacity across 50 states and ensures that a change in federal administration cannot eliminate enforcement.

## The Penalty Structure

The original proposal set a flat $100 million cap. Our analysts identified the problem: $100 million is less than 0.025% of Alphabet's annual revenue. It's a parking ticket.

We adopt the GDPR model: **the greater of $100 million OR 4% of global annual revenue**, whichever is larger. For Alphabet, that's up to $16 billion maximum exposure. For a smaller company, the $100 million floor still applies.

**Age-tiered penalty escalation:** Penalties are larger for violations involving younger children.

| **Violation** | **Under-13 Penalty** | **Ages 13-17 Penalty** |
|---|---|---|
| Missing required parental control features | $10M per feature per quarter | $5M per feature per quarter |
| Minor accessing high-risk feature without required consent | $20M per verified violation | $10M per verified violation |
| Failure to send a crisis alert | $20M per incident | $10M per incident |
| Data breach without notification | $50M per incident | $50M per incident |
| Willful non-compliance (FTC finding) | 4× all penalties + federal fund clawback | 2× all penalties + federal fund clawback |
| **Maximum per year** | **Greater of $200M or 8% of global revenue** | **Greater of $100M or 4% of global revenue** |

**Federal fund clawback** remains the most powerful tool for large AI companies competing for government contracts. Companies receiving federal AI research grants, defense contracts, or education technology funding lose those funds upon willful non-compliance. For companies like OpenAI and Anthropic, losing federal contracts is a more powerful deterrent than any fine.

## Annual Audits and Transparency

Every covered company must submit an annual public compliance report to the FTC including:

- Percentage of minor accounts by age tier with appropriate consent on file
- Crisis alert volume, response time, and detection accuracy
- Content control adoption rates by age tier

- Any incidents involving harm to minors and company response

These reports are public. Transparency creates accountability. Reporters, researchers, and parents will read them.

---

## Part 7: Timeline

Eighteen months from enactment to full compliance. This matches the COPPA 2013 rule update — the most recent comparable regulation. OpenAI built parental controls in a single product cycle. Google Family Link and Microsoft Family Safety already exist. Companies with $100B+ in revenue and tens of thousands of engineers can extend existing safety infrastructure in 18 months.

| Timeframe | Milestone |
|---|---|
| **Months 1–6** | FTC publishes technical standards; covered companies begin registration and fee payment |
| **Months 6–12** | Covered companies submit compliance plans to FTC; age-tiered systems in development |
| **Months 12–18** | All five features live for all verified minor accounts, with age-tiered protections active |
| **Annually** | FTC receives compliance reports; publishes public transparency summary; private actions can proceed |

---

## Five Specific Asks

### For Federal Education Leaders

Champion the PROTECT Act as part of the National AI Policy Framework's child safety priority — which the White House has already declared its first recommendation. Connect parental control requirements to Title IV technology funding so AI safety becomes part of what schools expect from every AI tool they recommend to students and families.

### For State Governors

Issue executive guidance directing state education agencies to prioritize AI platforms that already meet PROTECT Act standards when making purchasing decisions. The federal law is coming — position your state ahead of it.

### For Congress

Move swiftly to introduce and pass the PROTECT Act. The White House asked for action on March 20, 2026. Oregon passed a state version unanimously on March 5. KOSA passed the Senate 91-3. Child safety is not a partisan issue — it is the one issue where everyone agrees. Use the bright-line statutory triggers we define, not vague agency delegation. Use the GDPR penalty model so fines actually deter behavior. Fund enforcement through industry registration fees, not FTC budget battles.

### For the AI Industry

Don't wait for the law. Build these tools now. OpenAI showed in September 2025 that it can be done — and then went further, joining Common Sense Media to publicly support mandatory controls in January 2026. That is what industry leadership looks like. The companies that lead on parental controls will have a durable competitive advantage with families. The ones that don't will face the full weight of this legislation, private litigation, and the reputational cost of being the company that let three children die without building the tools that could have helped.

### For All Policymakers

Treat AI child safety the way you treat car safety, food safety, and drug safety. You don't let unsafe cars on the road. You don't put untested drugs on shelves. You required childproof caps on aspirin after children were poisoned. AI chatbots are being accessed by children as young as 8 — with demonstrably dangerous features — without any of these protections. Three children are dead. The technology to protect them exists. We know how to build it. What does not exist is the law requiring every company to deploy it.

That is what the PROTECT Act is. It is time to pass it.

---

## About This Proposal

This proposal was created by **Madhesh Saravanan**, a student and member of **Team AI Glitch Busters** — competing in the Presidential AI Challenge 2026, Track II Elementary. The team includes **Aariv Kumar**, **Madhesh Saravanan**, and **Mishika Chauhan**.

Madhesh's interest in this issue came from his own experience as a student using AI tools and noticing that his parents had no visibility into those interactions. He started asking: why don't AI companies have the same kinds of parental controls that video games and TVs already have? The research confirmed his instinct — and revealed the harm was far more serious than he initially realized.

This proposal is one of three complementary proposals from Team AI Glitch Busters — a complete framework for safer AI in American education.

---

## Appendix A: Key Definitions

**AI Chatbot:** A computer program powered by artificial intelligence that has conversations with people using natural language. Examples include ChatGPT, Gemini, Claude, Meta AI, and Character.AI.

**Parental Dashboard:** A secure online portal or app where a parent or guardian can see summaries of their child's AI activity, set content controls, adjust usage limits, and receive alerts. Opt-in; not activated by default.

**Crisis Alert:** An immediate notification sent to a parent or guardian when an AI detects language in a child's conversation related to self-harm, suicide, or personal emergency. Mandatory for all ages, all services, always active. The only non-opt-in feature.

**High-Risk AI Features:** AI features with documented potential for serious harm to minors, specifically: companion chatbots (persistent persona relationships), emotional support chatbots (AI positioned as therapist/romantic partner), roleplay chatbots (unrestricted character creation), and adult content features. These require consent as determined by the age-tiered approach.

**Covered Company:** An AI company required to comply with the PROTECT Act. Coverage is triggered by: (1) more than $1 billion in annual revenue or federal AI funding, (2) meeting all three bright-line statutory thresholds (15%+ minor users + companion conversations + documented incident), or (3) any link to a documented minor fatality or federal enforcement action involving minors.

**Age-Tiered Approach:** The PROTECT Act's three-tier system that calibrates requirements to developmental stage: Under 13 (full parental consent for all AI); Ages 13-15 (parental consent for high-risk features, safe defaults elsewhere); Ages 16-17 (teen self-consent for high-risk features, opt-in parental summaries). Crisis alerts mandatory for all three tiers.

**Deployment-Based Obligation:** The PROTECT Act's rule that compliance attaches to whoever deploys an AI service to end users — regardless of whether the underlying model is open-source or proprietary. The deployer is responsible; the model's license is irrelevant.

**FTC Child AI Safety Division:** A dedicated enforcement unit within the Federal Trade Commission, funded by covered companies' annual registration fees and earmarked penalty revenue, with specific statutory authority to enforce the PROTECT Act.

---

## Appendix B: Comparison Table

| Feature | COPPA (1998/2025) | KOSA (Proposed) | GUARD Act (Proposed) | EU DSA Article 28 | PROTECT Act v2 |
|---|---|---|---|---|---|
| **Covers AI chatbots specifically?** | ❌ No | ⚠️ Partially | ✅ Yes (bans them) | ⚠️ Partially | ✅ Yes |
| **Covers ages 13–17?** | ❌ No (under-13 only) | ✅ Yes | ✅ Yes | ✅ Yes | ✅ Yes (age-tiered) |
| **Requires parental dashboard?** | ❌ No | ❌ No | ❌ No | ⚠️ Recommends | ✅ Opt-in, mandated |
| **Requires crisis alerts?** | ❌ No | ❌ No | ❌ No | ❌ No | ✅ Mandatory always-on |
| **Requires content controls?** | ❌ No | ⚠️ Partial | N/A (bans service) | ⚠️ Recommends | ✅ Default-safe, adjustable |
| **Consent gate approach?** | ✅ Under-13 only | ❌ No | ✅ All AI use (ban) | ⚠️ Recommends | ✅ High-risk features only, age-tiered |
| **First Amendment compliant?** | ✅ Yes | ⚠️ Uncertain | ❌ Likely not | N/A (EU) | ✅ Yes (narrowly tailored) |

| Feature | COPPA (1998/2025) | KOSA (Proposed) | GUARD Act (Proposed) | EU DSA Article 28 | PROTECT Act v2 |
|---|---|---|---|---|---|
| **Preserves beneficial AI access?** | ✅ Yes | ✅ Yes | ❌ No (outright ban) | ✅ Yes | ✅ Yes |
| **Revenue-scaled penalties?** | ❌ Flat cap | N/A | N/A | ✅ 4% global revenue | ✅ GDPR model + age multiplier |
| **Covers Character.AI?** | ❌ No | ❌ No | ✅ Yes (by banning) | N/A (EU) | ✅ Via bright-line triggers |
| **Open-source loophole?** | N/A | N/A | N/A | ✅ Closed | ✅ Deployment-based, closed |
| **Dedicated enforcement funding?** | ❌ No | N/A | N/A | ✅ Partially | ✅ Registration fees + earmarked penalties |
| **Age-tiered teen privacy?** | ✅ Under-13 only | ❌ No | ❌ No | ⚠️ Partial | ✅ Three-tier approach |
| **Private right of action?** | ✅ Yes | ✅ Yes | ✅ Yes | ⚠️ Limited | ✅ Yes, strengthened |

---

# Appendix C: Sources

**Original Sources (Version 1.0)**


1. Pew Research Center. "What Parents Say About Their Teen's AI Use." February 24, 2026. pewresearch.org
2. Pew Research Center. "How Teens Use and View AI." February 24, 2026. pewresearch.org

3. Pew Research Center. "Teens, Social Media and AI Chatbots 2025." December 2025. pewresearch.org
4. National Center for Education Statistics (NCES). Common Core of Data, 2022–23. nces.ed.gov
5. National Parents Union / Echelon Insights. Survey of 1,511 K-12 public school parents on AI guardrails. February 12–18, 2026. nationalparentsunion.org
6. EdWeek. "How Do Parents Want Schools to Handle AI?" March 2026. edweek.org
7. New York Times. "Google and Character.AI to Settle Lawsuit Over Teenager's Death." January 7, 2026.
8. CBS News. "AI company, Google settle lawsuit over Florida teen's suicide linked to Character.AI chatbot." January 8, 2026.
9. Stanford Medicine. "Why AI Companions and Young People Can Make for a Dangerous Mix." August 27, 2025.
10. Common Sense Media / Stanford Brainstorm Lab. "Major AI Chatbots Unsafe for Teen Mental Health Support." 2025.
11. Common Sense Media. "The Dawn of the AI Era." September 18, 2024.
12. JAMA. "Mitigating Suicide Risk for Minors Involving AI Chatbots." January 27, 2026.
13. CDC. "2023 Youth Risk Behavior Survey Results." Published October 2024. cdc.gov
14. U.S. Surgeon General. "Social Media and Youth Mental Health Advisory." May 2023. hhs.gov
15. FTC Press Release. "Finalizes Changes to Children's Privacy Rule (COPPA 2.0)." January 16, 2025. ftc.gov
16. FTC Press Release. "Launches Inquiry into AI Chatbots Acting as Companions." September 11, 2025. ftc.gov
17. FTC v. Epic Games. $520 million settlement. December 2022. ftc.gov
18. FTC v. Google/YouTube. $170 million COPPA penalty. 2019. ftc.gov
19. OpenAI. "Introducing Parental Controls." September 29, 2025. openai.com
20. The Verge. "Character.AI Has Retrained Its Chatbots to Stop Chatting Up Teens." December 12, 2024.
21. CNN. "Character.AI Will No Longer Let Teens Chat With Its Chatbots." October 29, 2025.
22. NPR. "Character.AI Chatbot Suggested Child Should Harm Parents Over Screen Time." December 2024.
23. Sen. Josh Hawley / Sen. Chris Murphy. GUARD Act introduction. October 28, 2025. hawley.senate.gov
24. Reuters. "European Lawmakers Seek EU-Wide Minimum Age to Access AI Chatbots." November 26, 2025.
25. EU Commission. "Guidelines on the Protection of Minors Under the Digital Services Act (DSA)." 2025.
26. MultiState. "Eight States Enact Minor Social Media Bans." October 8, 2025.
27. Kaiser Family Foundation. V-chip effectiveness survey. 2007.
28. Fortune Business Insights. "Parental Control Software Market." 2025.
29. Yale / MEPS. "Child Mental Health Services Cost $31 Billion Per Year." 2021.
30. Yale News. "U.S. Mental Illness Costs $282 Billion Annually." 2024.

31. College Board. AI in education survey (69% of high school students use AI for schoolwork). May 2025.
32. AI CERTs. "Therapy Chatbots and the Mental Health Crisis: Urgent Reforms." March 2026.
33. Cyberbullying Research Center. "How Platforms Should Build AI Chatbots for Youth Safety." January 14, 2025.
34. Virginia Tech News. "Experts Say Parental Controls Are Good, But AI Still Needs More Oversight." September 17, 2025.
35. Programs.com. "The Latest AI in Education Statistics (2026)." February 11, 2026.
36. Ireland Data Protection Commission v. TikTok. €345M GDPR fine. 2023.

**New Sources Added in Version 2.0**

37. TruLaw / Court filings. Adam Raine and Juliana Peralta deaths linked to Character.AI. 2025-2026.
38. Center for Countering Digital Hate (CCDH). "Lethal Listening: AI Chatbots and the Facilitation of Violence." March 11, 2026. counterhate.com
39. Axios / Court documents. Meta internal chatbot safety testing results (66.8% / 63.6% / 54.8% failure rates). February 2026.
40. White House. "National AI Policy Framework." March 20, 2026. whitehouse.gov
41. Oregon Legislature. Senate Bill 1546. Passed March 5, 2026. oregonlegislature.gov
42. Common Sense Media / OpenAI. Joint announcement supporting mandated AI child safety controls. January 9, 2026.
43. eSafety Commissioner (Australia). Age-Restricted Material Codes. Effective March 9, 2026. esafety.gov.au
44. UK Government / Ofcom. Online Safety Act enforcement updates. 2025-2026.
45. CPSC / EBSCO Research. Poison Prevention Packaging Act: 87% reduction in child poisoning deaths, 1972-2000.
46. EAST (Eastern Association for the Surgery of Trauma). Child car seat mandates: 58-70% injury reduction evidence review. Ongoing.
47. Columbia Business School. "U.S. Mental Illness Costs $282 Billion Per Year." 2024.
48. Judge Conway. Character.AI wrongful death claims can proceed. Court ruling. May 2025.
49. Children and Screens. "Parental Monitoring, Trust, and Teen Online Safety." 2025.
50. Journal of Children and Media. Rapid evidence review on parental monitoring software and teen outcomes. 2023.
51. Internet Watch Foundation. AI-generated CSAM report: 26,385% increase in video content. 2026.
52. National Highway Traffic Safety Administration. Child traffic fatality reduction data, 1970s-present.
53. West Virginia v. EPA. 597 U.S. 697 (2022). Major questions doctrine.
54. Free Speech Coalition v. Paxton. Supreme Court (2025). Age verification for obscene content.

55. NetChoice court decisions. Federal court injunctions on state minor social media access laws. 2024-2025.
56. Orrick analysis. EU AI Act open-source exemption scope. September 2024.

---

**Document End**

*PROTECT Act - The Parent's Right to Oversee Technology and Ensure Child-Safe AI Transparency Act*
*Presidential AI Challenge 2025–2026*
*Team AI Glitch Busters: Aariv Kumar, Madesh Saravanan, Mishika Chauhan*
*Supervisor: Ashwin Kumar, Wiz Kidz*
*Contact: https://garage.wizkidz.ai/ai-glitch-buster/*